

Group Final project: Baseline Implementation and Investigation

8-Hengke Jiang, 13-Yuting Jiang, 20-Rock, 30-Ethan, 35-Jade Lin

Table of Contents

1 Text classification and dataset

2 BERT and modified BERT

3 Implement and result

4 Further steps to be taken

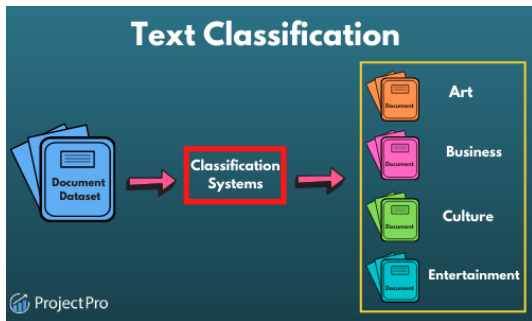
Text classification

Why Text classification is still alive?

Text classification applications

- Customer Support and Feedback Analysis: Automatically classifying customer queries into categories like technical support, billing issues, product inquiries, etc;
- Content Moderation and Filtering: Filtering and moderating user-generated content on social media platforms, forums, and online communities;
- Information Retrieval: Improving search engine results by categorizing and organizing documents or web pages based on their topics or relevance;
- Filtering incoming emails to separate spam or junk emails from legitimate messages;

Text classification



- News Categorization: Organizing news articles and updates into different categories or topics for readers.

Dataset and SOA

- 120,000 training articles, 7600 testing articles;
- Categories divided: World, Sports, Business, Science/Technology
- Format: Each article is represented as a text string along with its corresponding label indicating the topic category.




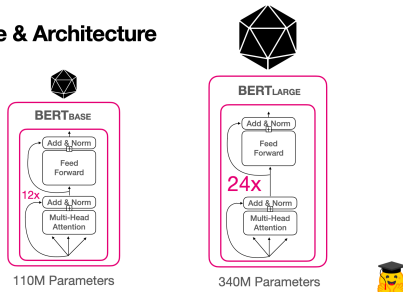
Table of Contents

- 1 Text classification and dataset
- 2 BERT and modified BERT**
- 3 Implement and result
- 4 Further steps to be taken

BERT

Base on encoder part of transformer[5], BERT(Bidirectional Encoder Representations from Transformers)[1] is developed to do tasks as: Text classification, Question Answering, Text Generation, Language Understanding, and Sentence Embeddings.

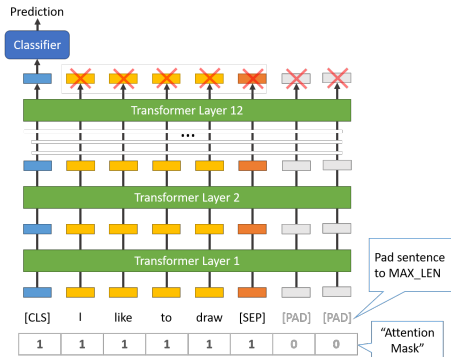
BERT Size & Architecture



Attention: without decoder part, BERT is not used for translation.

Feature of BERT-base

- 30,000 token vocabulary
- Transformer Blocks: 12 layers;
- Hidden Size: 768 (number of units in the hidden layers);
- Attention Heads: 12
- Number of parameters: 110M




Modified BERT: ALBERT

ALBERT: a lite BERT for self-supervised learning of language representations[4].



Motivation: at some point further model increases become harder due to GPU/TPU memory limitations and longer training times. Improve BERT in 2 ways with the same number of parameters(110M).

- Tech 1: Decomposing the large vocabulary embedding matrix into two small matrices to separate the size of the hidden layers from the size of vocabulary embedding.

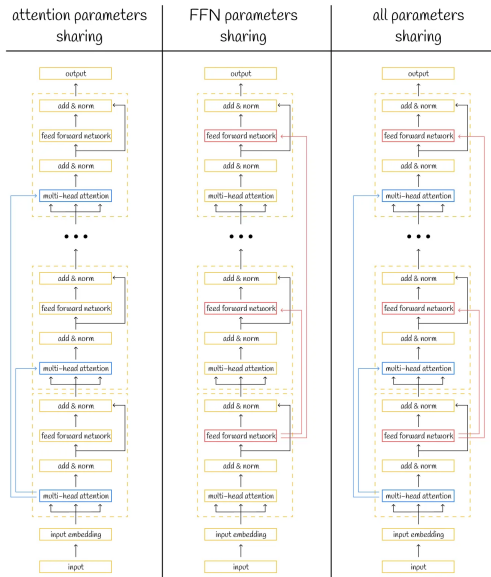


Modified BERT: ALBERT

- Tech 2: cross-layer parameter sharing prevents the parameter from growing with the depth of the network;



Modified BERT: ALBERT



Modified BERT: DEBERTA

DeBERTa: Decoding-enhanced BERT with Disentangled Attention[3]

Motivation: Improvement with the number of parameters close to BERT(89M vs. 110M) in two ways.

- 1. Disentangled attention: Each word is represented using two vectors that encode its content and position, respectively, and the attention weights among words are computed using disentangled matrices based on their contents and relative positions, respectively.

For a token at position i in a sequence, we represent it using two vectors, $\{\mathbf{H}_i\}$ and $\{\mathbf{P}_{i|j}\}$, which represent its content and relative position with the token at position j , respectively. The calculation of the cross attention score between tokens i and j can be decomposed into four components as

$$\begin{aligned} A_{i,j} &= \{\mathbf{H}_i, \mathbf{P}_{i|j}\} \times \{\mathbf{H}_j, \mathbf{P}_{j|i}\}^\top \\ &= \mathbf{H}_i \mathbf{H}_j^\top + \mathbf{H}_i \mathbf{P}_{j|i}^\top + \mathbf{P}_{i|j} \mathbf{H}_j^\top + \mathbf{P}_{i|j} \mathbf{P}_{j|i}^\top \end{aligned} \quad (2)$$

- 2. Enhanced mask decoder: Each word in DeBERTa is represented using two vectors that encode its content and position, respectively, and the attention weights among words are computed using disentangled matrices based on their contents and relative positions, respectively.

Table of Contents

- 1 Text classification and dataset
- 2 BERT and modified BERT
- 3 Implement and result**
- 4 Further steps to be taken




parameter and devices(pretrained)




- number of epoches = 6, batchsize = 32;
- divide training set into 2 part and shuffles, each epoch cost 4~8 min;
- optimizer: AdamW, with linearly decreasing learning rate($6e-5 \sim 1e-5$)
- 3 models implement.




```
100% | 1875/1875 [04:44<01:00, 6.581t/s]
Epoch 1/6 completed. Train Loss: 0.2283, Train Accuracy: 0.9185, Val Loss: 0.1928, Val Accuracy: 0.9347, F1 Score: 0.9184
100% | 1875/1875 [04:44<01:00, 6.581t/s]
Epoch 2/6 completed. Train Loss: 0.1338, Train Accuracy: 0.9545, Val Loss: 0.1778, Val Accuracy: 0.9488, F1 Score: 0.9545
100% | 1875/1875 [04:45<00:00, 6.581t/s]
Epoch 3/6 completed. Train Loss: 0.0840, Train Accuracy: 0.9705, Val Loss: 0.1974, Val Accuracy: 0.9432, F1 Score: 0.9705
100% | 1875/1875 [04:44<01:00, 6.581t/s]
Epoch 4/6 completed. Train Loss: 0.0448, Train Accuracy: 0.9842, Val Loss: 0.2315, Val Accuracy: 0.9411, F1 Score: 0.9843
100% | 1875/1875 [04:44<01:00, 6.581t/s]
Epoch 5/6 completed. Train Loss: 0.0235, Train Accuracy: 0.9918, Val Loss: 0.2585, Val Accuracy: 0.9426, F1 Score: 0.9918
100% | 1875/1875 [04:44<00:00, 6.581t/s]
Epoch 6/6 completed. Train Loss: 0.0103, Train Accuracy: 0.9965, Val Loss: 0.2980, Val Accuracy: 0.9447, F1 Score: 0.9965
Test Accuracy: 0.9447
Test F1 Score: 0.9447
```

```
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 1/6 completed. Train Loss: 0.5346, Train Accuracy: 0.7884, Val Loss: 1.3883, Val Accuracy: 0.2576, F1 Score: 0.7877
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 2/6 completed. Train Loss: 0.7819, Train Accuracy: 0.6556, Val Loss: 0.3580, Val Accuracy: 0.8950, F1 Score: 0.6554
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 3/6 completed. Train Loss: 0.2935, Train Accuracy: 0.9129, Val Loss: 0.3441, Val Accuracy: 0.8955, F1 Score: 0.9127
100% | 1875/1875 [05:20<00:00, 5.841t/s]
Epoch 4/6 completed. Train Loss: 0.2123, Train Accuracy: 0.9287, Val Loss: 0.2627, Val Accuracy: 0.9134, F1 Score: 0.9286
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 5/6 completed. Train Loss: 0.1744, Train Accuracy: 0.9452, Val Loss: 0.2694, Val Accuracy: 0.9176, F1 Score: 0.9452
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 6/6 completed. Train Loss: 0.1392, Train Accuracy: 0.9564, Val Loss: 0.2554, Val Accuracy: 0.9221, F1 Score: 0.9564
Test Accuracy: 0.9221
Test F1 Score: 0.9222
```

```
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 1/6 completed. Train Loss: 0.2517, Train Accuracy: 0.9151, Val Loss: 0.1908, Val Accuracy: 0.9361, F1 Score: 0.9150
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 2/6 completed. Train Loss: 0.1477, Train Accuracy: 0.9493, Val Loss: 0.1797, Val Accuracy: 0.9418, F1 Score: 0.9493
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 3/6 completed. Train Loss: 0.0985, Train Accuracy: 0.9657, Val Loss: 0.1938, Val Accuracy: 0.9408, F1 Score: 0.9657
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 4/6 completed. Train Loss: 0.0634, Train Accuracy: 0.9786, Val Loss: 0.2148, Val Accuracy: 0.9429, F1 Score: 0.9786
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 5/6 completed. Train Loss: 0.0348, Train Accuracy: 0.9878, Val Loss: 0.2342, Val Accuracy: 0.9413, F1 Score: 0.9878
100% | 1875/1875 [06:19<00:00, 4.951t/s]
Epoch 6/6 completed. Train Loss: 0.0192, Train Accuracy: 0.9937, Val Loss: 0.2824, Val Accuracy: 0.9418, F1 Score: 0.9937
Test Accuracy: 0.9418
Test F1 Score: 0.9416
```





Summarazie: BERT > ALBERT > DeBERTa, partly because testing part is divided along with training part.

```
100% | 1875/1875 [04:44<00:00, 6.581t/s]
Epoch 1/6 completed. Train Loss: 0.2393, Train Accuracy: 0.9185, Val Loss: 0.1920, Val Accuracy: 0.9347, F1 Score: 0.9184
100% | 1875/1875 [04:44<00:00, 6.581t/s]
Epoch 2/6 completed. Train Loss: 0.1339, Train Accuracy: 0.9545, Val Loss: 0.1778, Val Accuracy: 0.9480, F1 Score: 0.9545
100% | 1875/1875 [04:45<00:00, 6.581t/s]
Epoch 3/6 completed. Train Loss: 0.0840, Train Accuracy: 0.9764, Val Loss: 0.1974, Val Accuracy: 0.9432, F1 Score: 0.9766
100% | 1875/1875 [04:44<00:00, 6.581t/s]
Epoch 4/6 completed. Train Loss: 0.0448, Train Accuracy: 0.9842, Val Loss: 0.2335, Val Accuracy: 0.9411, F1 Score: 0.9843
100% | 1875/1875 [04:44<00:00, 6.581t/s]
Epoch 5/6 completed. Train Loss: 0.0239, Train Accuracy: 0.9918, Val Loss: 0.2585, Val Accuracy: 0.9426, F1 Score: 0.9918
100% | 1875/1875 [04:44<00:00, 6.581t/s]
Epoch 6/6 completed. Train Loss: 0.0102, Train Accuracy: 0.9965, Val Loss: 0.2980, Val Accuracy: 0.9447, F1 Score: 0.9965
Test Accuracy: 0.9447
Test F1 Score: 0.9447
```

```
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 1/6 completed. Train Loss: 0.5046, Train Accuracy: 0.7884, Val Loss: 1.3883, Val Accuracy: 0.2576, F1 Score: 0.7877
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 2/6 completed. Train Loss: 0.7819, Train Accuracy: 0.6556, Val Loss: 0.3580, Val Accuracy: 0.8900, F1 Score: 0.6554
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 3/6 completed. Train Loss: 0.2809, Train Accuracy: 0.9129, Val Loss: 0.3441, Val Accuracy: 0.8995, F1 Score: 0.9127
100% | 1875/1875 [05:20<00:00, 5.841t/s]
Epoch 4/6 completed. Train Loss: 0.2222, Train Accuracy: 0.9287, Val Loss: 0.2622, Val Accuracy: 0.9134, F1 Score: 0.9286
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 5/6 completed. Train Loss: 0.1744, Train Accuracy: 0.9452, Val Loss: 0.2694, Val Accuracy: 0.9176, F1 Score: 0.9452
100% | 1875/1875 [05:21<00:00, 5.841t/s]
Epoch 6/6 completed. Train Loss: 0.1392, Train Accuracy: 0.9594, Val Loss: 0.2554, Val Accuracy: 0.9221, F1 Score: 0.9594
Test Accuracy: 0.9221
Test F1 Score: 0.9222
```

```
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 1/6 completed. Train Loss: 0.2217, Train Accuracy: 0.9151, Val Loss: 0.1909, Val Accuracy: 0.9361, F1 Score: 0.9150
100% | 1875/1875 [06:19<00:00, 4.951t/s]
Epoch 2/6 completed. Train Loss: 0.1477, Train Accuracy: 0.9493, Val Loss: 0.1737, Val Accuracy: 0.9418, F1 Score: 0.9493
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 3/6 completed. Train Loss: 0.0866, Train Accuracy: 0.9647, Val Loss: 0.1938, Val Accuracy: 0.9408, F1 Score: 0.9647
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 4/6 completed. Train Loss: 0.0634, Train Accuracy: 0.9788, Val Loss: 0.2149, Val Accuracy: 0.9429, F1 Score: 0.9788
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 5/6 completed. Train Loss: 0.0368, Train Accuracy: 0.9876, Val Loss: 0.2342, Val Accuracy: 0.9413, F1 Score: 0.9876
100% | 1875/1875 [06:19<00:00, 4.941t/s]
Epoch 6/6 completed. Train Loss: 0.0192, Train Accuracy: 0.9937, Val Loss: 0.2924, Val Accuracy: 0.9418, F1 Score: 0.9937
Test Accuracy: 0.9418
Test F1 Score: 0.9416
```

R8

The R8 dataset is a subset of the Reuters-21578 dataset, which is commonly used for text classification tasks in natural language processing (NLP). The R8 dataset consists of eight categories of news documents

Categories: The dataset is divided into the following eight categories:

Acquisitions Crude Earn Grain Interest Money-fx Ship Trade

The dataset is typically split into a training set and a testing set, allowing for the evaluation of machine learning models. The standard split often used is 5,485 documents for training and 2,189 documents for testing.

- smaller than AG News, affordable;

Implement without pretrain

- number of epoch = 10, batchsize = 4(for time saving)

```
class BERT(nn.Module):
    def __init__(self, vocab_size, hidden=768, n_layers=12, attn_heads=12, dropout=0.1):

        super().__init__()
        self.hidden = hidden
        self.n_layers = n_layers
        self.attn_heads = attn_heads

        # paper noted they used 4*hidden_size for ff_network_hidden_size
        self.feed_forward_hidden = hidden * 4

        # embedding for BERT, sum of positional, segment, token embeddings
        self.embedding = BERTEmbedding(vocab_size=vocab_size, embed_size=hidden)

        # multi-layers transformer blocks, deep network
        self.transformer_blocks = nn.ModuleList(
            [TransformerBlock(hidden, attn_heads, hidden * 4, dropout) for _ in range(n_layers)])

    def forward(self, x, segment_info):
        # attention masking for padded token
        # torch.ByteTensor([batch_size, 1, seq_len, seq_len])
        mask = (x > 0).unsqueeze(1).repeat(1, x.size(1), 1).unsqueeze(1)

        # embedding the indexed sequence to sequence of vectors
        x = self.embedding(x, segment_info)

        # running over multiple transformer blocks
        for transformer in self.transformer_blocks:
            x = transformer.forward(x, mask)

        return x
```






Table of Contents

- 1 Text classification and dataset
- 2 BERT and modified BERT
- 3 Implement and result
- 4 Further steps to be taken

Old tech: change the parameter used, like learning rate and batchsize;
New tech: DEBERTA V3[2]

- 1 improves the original DeBERTa model by replacing mask language modeling (MLM) with replaced token detection (RTD), a more sample-efficient pre-training task;
- 2 not yet updated in text classification task.



- [1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova.
Bert: Pre-training of deep bidirectional transformers for language understanding, 2019.
- [2] Pengcheng He, Jianfeng Gao, and Weizhu Chen.
Debertav3: Improving deberta using electra-style pre-training with gradient-disentangled embedding sharing, 2023.
- [3] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen.
Deberta: Decoding-enhanced bert with disentangled attention, 2021.
- [4] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut.
Albert: A lite bert for self-supervised learning of language representations, 2020.

- [5] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin.
Attention is all you need, 2023.